



Overview and Contributions

- A differentiable implementation of every direct form digital filter for PyTorch, utilising parallel algorithms to accelerate both the forward pass and gradient computation.
- Analytical gradients are derived for any reverse-mode automatic differentiation frameworks.
- Over 1000x speedup compared to naive implementations.

Problem Statement: Recursive Filters Are Slow For AD

An M^{th} -order time-invariant linear system has the following direct form (DF) difference equation:

$$y(n) = b_0x(n) + b_1x(n-1) + \dots + b_Mx(n-M) - a_1y(n-1) - a_2y(n-2) - \dots - a_My(n-M) \quad (1)$$

where b_i, a_i are the feed-forward and feedback coefficients, $x(n)$ and $y(n)$ are the input and output signals. The definition of $y(n)$ is recursive.

An implementation of $\mathbf{v}(n+1) = \mathbf{Av}(n) + \mathbf{z}(n)$ in PyTorch:

```
v = []
vn = torch.zeros(M)
for n in range(N):
    vn = torch.addmv(z[n], A, vn)
    v.append(vn)
v = torch.stack(v)
```

- Runtime: $N \times (\text{torch.* overhead} + \text{execution time})$
- `torch.compile`: Bad at optimising for-loops
- Dedicated recursion kernel has a runtime of just `torch.* overhead + execution time`
 - Support ✓: JAX and TensorFlow
 - WIP ✗: PyTorch [1]

Previous Work: Frequency-Sampling

Filtering in the frequency domain only needs fixed number of calls to `torch.fft.*`:

$$H(e^{j\omega}) = \frac{b_0 + b_1e^{-j\omega} + \dots + b_Me^{-jM\omega}}{1 + a_1e^{-j\omega} + \dots + a_Me^{-jM\omega}}. \quad (2)$$

Issues commonly observed:

- Frequency-domain sampling → time-domain aliasing
- $Y(j\omega) = H(j\omega)X(j\omega)$ results in circular convolution artefacts
- Less efficient for low-order filters

Proposed Method: State-Space and Custom Extension

We implement filtering as a custom C++/CUDA extension, avoiding the accumulated overhead proportional to N . This has been used before [2, 3] but

- Only applies to a limited set of filters (DF-I/II and all-pole)
- No parallelisation across the time dimension
- Does not include gradients to the initial state

We rewrite Eq. (1) into state-space form:

$$\mathbf{v}(n+1) = \mathbf{Av}(n) + \mathbf{Bx}(n) \quad (3)$$

$$y(n) = \mathbf{C}^T \mathbf{v}(n) + \mathbf{Dx}(n) \quad (4)$$

$$\mathbf{A} = \begin{bmatrix} -\mathbf{a}^T \\ \mathbf{I} \mathbf{0} \end{bmatrix}, \mathbf{C} = \begin{bmatrix} b_1 - a_1b_0 \\ b_2 - a_2b_0 \\ \vdots \\ b_M - a_Mb_0 \end{bmatrix}, \mathbf{B} = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, D = b_0. \quad (5)$$

Benefits. State-space form covers transposed direct forms (TDF) (more numerically robust and common) with $\mathbf{A} \mapsto \mathbf{A}^T$, $\mathbf{C} \mapsto \mathbf{B}$, $\mathbf{B} \mapsto \mathbf{C}$. Eq. (3) can be parallelised by computing the following associative scan [4]:

$$(\mathbf{0}, \mathbf{v}(n)) = (\mathbf{0}, \mathbf{v}(0)) \oplus (\mathbf{A}, \mathbf{Bx}(0)) \oplus \dots \oplus (\mathbf{A}, \mathbf{Bx}(n-1)), \quad (6)$$

where $(\mathbf{A}, \mathbf{z}) \oplus (\mathbf{A}', \mathbf{z}') \mapsto (\mathbf{A}'\mathbf{A}, \mathbf{A}'\mathbf{z} + \mathbf{z}')$.

Gradient Backpropagation

Let $\mathcal{L} = f(y(0), y(1), \dots)$ be the scalar we want to minimise. Given the output gradients $\frac{\partial \mathcal{L}}{\partial y(n)}$, the gradients to the variables in the filter can be computed as:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{C}} = \sum_{n \geq 0} \frac{\partial \mathcal{L}}{\partial y(n)} \mathbf{v}(n)^T, \quad \frac{\partial \mathcal{L}}{\partial D} = \sum_{n \geq 0} \frac{\partial \mathcal{L}}{\partial y(n)} x(n), \quad (7)$$

$$\mathbf{u}(n) = \mathbf{A}^T \mathbf{u}(n+1) + \mathbf{C} \frac{\partial \mathcal{L}}{\partial y(n+1)}, \quad (8)$$

$$\frac{\partial \mathcal{L}}{\partial x(n)} = \mathbf{B}\mathbf{u}(n) + D \frac{\partial \mathcal{L}}{\partial y(n)}, \quad (9)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{v}(0)} = \mathbf{u}(-1)^T, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{B}} = \sum_{n \geq 0} \mathbf{u}(n)^T x(n), \quad \frac{\partial \mathcal{L}}{\partial \mathbf{A}} = \sum_{n \geq 0} \mathbf{v}(n) \mathbf{u}(n)^T. \quad (10)$$

Key Insight. Eq. (8) and (9) form a TDF filter in state-space form; similarly, the backpropagation of a TDF filter is a DF filter. Moreover, we can use the same parallel scan method to accelerate Eq. (8).

Experimental Results

We benchmarked the recursion Eq. (3) with $M = 2$ and sequence length ranging from 2^{14} to 2^{20} in PyTorch 2.8 on an Intel i7-7700K and an NVIDIA RTX 5060 Ti.

- Naive:** Simple for-loop implementation.
- RTF:** Frequency-domain filtering using Eq. (2).
- Unrolled:** Parallel scan using `torch.matmul` [5].
- EXT:** Our custom C++/CUDA extension.
- Diag-EXT:** Diagonalise $\mathbf{A} = \mathbf{V}\text{diag}(\lambda)\mathbf{V}^{-1}$ and perform scan in the eigenbasis [6].




Figure 1. Runtime of the recursion and its backpropagation with various signal lengths N and implementations.

- The naive implementation is remarkably slower and impractical.
- EXT consistently spends the least time computing gradients in most configurations.
- RTF is comparable to Diag-EXT in smaller N on the GPU. However, it is only slightly faster than Naive on the CPU.

Conclusions and Next Steps

- Our custom extension for direct form filters significantly reduces the runtime of both the forward pass and backpropagation compared to existing methods.
- We plan to extend the derivations to parameter-varying cases, a broader class of state-space models, and forward-mode differentiation. Analysing the numerical robustness of the analytical gradients is another interesting direction.

- [1] Wu et al. Control Flow Operators in PyTorch. 2025.
- [2] Forgiore and Piga. dynoNet: A neural network architecture for learning dynamical systems. en. *International Journal of Adaptive Control and Signal Processing* 35.4 (2021). ISSN: 1099-1115.
- [3] Yu and Fazekas. GOLF: A Singing Voice Synthesiser with Glottal Flow Wavetables and LPC Filters. *TISMIR* (Dec. 2024). doi: 10.5334/tismir.210.
- [4] Blelloch. Prefix Sums and Their Applications. Tech. rep. CMU-CS-90-190. School of Computer Science, Carnegie Mellon University, Nov. 1990.
- [5] Yu. Block-based Fast Differentiable IIR in PyTorch. <https://iamycy.github.io/posts/2025/06/28/unroll-ssm/>. June 2025.
- [6] Yu et al. DiffVox: A Differentiable Model for Capturing and Analysing Vocal Effects Distributions. 2025.