

PHILTORCH

*ACCELERATING AUTOMATIC DIFFERENTIATION
OF DIGITAL FILTERS IN PYTORCH*

CHIN-YUN YU

ADC²⁵
Bristol

About me

- Yu, Chin-Yun (尤靖允)
- PhD student in AI & Music at Queen Mary University of London
- Research interests
 - Music Information Retrieval
 - Voice synthesis
 - Audio Effects Modelling
 - Binaural Audio
- <https://iamycy.github.io/>



Sources

- Publications
 - “Singing voice synthesis using differentiable LPC and glottal-flow-inspired wavetables”, ISMIR 2023
 - “Differentiable all-pole filters for time-varying audio systems”, DAFx24
 - “DiffVox: A differentiable model for capturing and analysing vocal effects distributions”, DAFx25
- Blog posts
 - “Introduction to Differentiable Audio Synthesiser Programming”, ISMIR 2023 Tutorial
 - “Notes on Differentiable TDF-II Filter”, personal blog
 - “Block-based Fast Differentiable IIR in PyTorch”, personal blog

Agenda

- What are differentiable filters
- Efficiency problem in PyTorch
- Proposed solution: PhilTorch
- Benchmarks
- Tips for efficient usage

What Are Digital Filters?

- In the Z domain

$$H(z) = \frac{b_0 + b_1 z^{-1} + \cdots + b_M z^{-M}}{1 + a_1 z^{-1} + \cdots + a_M z^{-M}}$$

- In the time domain

$$\begin{aligned}y(n) = & b_0 x(n) + b_1 x(n - 1) + \cdots + b_M x(n - M) \\& - a_1 y(n - 1) - \cdots - a_M y(n - M)\end{aligned}$$

Usage

- Low/high/all-pass filters
- Shelving filters
- Peak filters
- Equalisers
- Subtractive synthesiser
- Phaser
- Linear prediction
- .etc



Differentiable Use Cases

- Differentiable Digital Signal Processing (DDSP)
 - Interpretable structure
 - strong inductive bias → highly efficient systems
- e.g.
 - Grey-box virtual analogue modelling
 - Synthesiser parameter retrieval




Fig. 3 from “Differentiable all-pole filters for time-varying audio systems”

Automatic Differentiation (AD) of Filters

- Differentiable everywhere? ✓
- Fast computation in PyTorch? ✗
- Problems
 - No native ops for recursions/autoregression

`torch.nn.functional.conv1d`

$$y(n) = b_0x(n) + b_1x(n - 1) + \dots + b_Mx(n - M)$$

$$-a_1y(n - 1) - \dots - a_My(n - M)$$

`torch.* ?`

Naïve AD of Recursions

```
h = zi
for xn in x.unbind(1):
    h = torch.addmm(xn, h, AT)
    results.append(h if out_idx is None else h[:, out_idx])
output = torch.stack(results, dim=1)
```

philtorch/lti/ssm.py, L127-131

runtime

torch.addmm

kernel

torch.addmm

kernel

...

torch.stack

kernel

F.conv1d

kernel

PhilTorch Φ

- A Python package for PyTorch
- Fast computation of linear discrete time filtering
- Same numerical robustness as SciPy
- Reverse- (and forward) mode AD
- Linear time-invariant filters
- Parameter-varying (time-varying) filters

```
pip install philtorch
```



Code



Things To Be Covered

The state-space formulation

Custom Ops for recursion

Backpropagation

Parallelisation

The State-Space Model (SSM)

... of a direct form (DF) II filter.


$$\mathbf{v}(n + 1) = \mathbf{Av}(n) + \mathbf{Bx}(n)$$

$$y(n) = \mathbf{C}^T \mathbf{v}(n) + D x(n)$$


$$\mathbf{A} = \begin{bmatrix} -\mathbf{a} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}, \quad \mathbf{B} = [1 \quad \dots \quad 0]$$
$$\mathbf{C} = [b_1 - b_0 a_1 \quad \dots \quad b_M - b_0 a_M], \quad D = b_0$$

Different Direct Forms

- DF-II
 - Used in TorchAudio
- Transposed DF-II (TDF-II)
 - Used in SciPy




$$H(z) = D + \mathcal{C}^T (zI - A)^{-1} \mathcal{B}$$



$$H(z) = D + \mathcal{B}^T (zI - A^T)^{-1} \mathcal{C}$$

Image source: Julius Smith , “Introduction to Digital Filters with Audio Applications”.

Writing Custom Operators in PyTorch



Output: Green
Input: Orange

$$\begin{aligned}\mathbf{v}(n+1) \\= \mathbf{A}\mathbf{v}(n) + \mathbf{B}\mathbf{x}(n) \\= \mathbf{A}\mathbf{v}(n) + \mathbf{u}(n)\end{aligned}$$

Backpropagation

...is a reverse-time recurrence with transposed \mathbf{A} .

- Gradients for the input

$$\frac{\partial \mathcal{L}}{\partial \mathbf{u}(n)} = \left(\mathbf{A}^T \frac{\partial \mathcal{L}}{\partial \mathbf{u}(n+1)}^\top + \frac{\partial \mathcal{L}}{\partial \mathbf{v}(n+1)}^\top \right)^\top$$

- Gradients for coefficients & initial condition

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}} = \sum_n \mathbf{v}(n) \frac{\partial \mathcal{L}}{\partial \mathbf{u}(n)}, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{v}(0)} = \frac{\partial \mathcal{L}}{\partial \mathbf{u}(0)} \mathbf{A}$$

Parallel implementation

Non-parallelisable and sequential ☹

$$\mathbf{v}(n+1) = \mathbf{A}(\mathbf{A}(\mathbf{A}(\dots) + \mathbf{u}(n-2)) + \mathbf{u}(n-1)) + \mathbf{u}(n)$$

Parallelisable ☺

- Parallel associative scan

$$(\mathbf{U}, \mathbf{x}) \oplus (\mathbf{V}, \mathbf{z}) \mapsto (\mathbf{V}\mathbf{U}, \mathbf{V}\mathbf{x} + \mathbf{z})$$

$$(\mathbf{0}, \mathbf{v}(n+1)) = (\mathbf{0}, \mathbf{v}(0)) \oplus (\mathbf{A}, \mathbf{u}(0)) \oplus (\mathbf{A}, \mathbf{u}(1)) \oplus \dots \oplus (\mathbf{A}, \mathbf{u}(n))$$

Diagonalised SSM

$$\tilde{\mathbf{v}}(n+1) = \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_M \end{bmatrix} \tilde{\mathbf{v}}(n) + \mathbf{P}^{-1} \mathbf{u}(n)$$

$$\mathbf{v}(n) = \mathbf{P}\tilde{\mathbf{v}}(n), \quad \mathbf{A} = \mathbf{P} \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_M \end{bmatrix} \mathbf{P}^{-1}$$

- Matrix multiplication → elementwise multiplications
- Equals to M first-order filters run in parallel

Alternative Solution w/o Custom Ops

- Sample unrolling
 - Andres Viso, “Implementing real-time Parallel DSP on GPUs”, ADC 2022.
- Pros: very efficient using `torch.matmul`
- Cons: still sequential

$$\begin{bmatrix} \mathbf{v}(n+1) \\ \vdots \\ \mathbf{v}(n+N-1) \\ \mathbf{v}(n+N) \end{bmatrix} = \begin{bmatrix} \mathbf{A} \\ \vdots \\ \mathbf{A}^{N-1} \\ \mathbf{A}^N \end{bmatrix} \mathbf{v}(n) + \begin{bmatrix} \mathbf{I} & \mathbf{0} & \cdots & \mathbf{0} \\ \mathbf{A} & \mathbf{I} & \cdots & \mathbf{0} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{A}^{N-1} & \mathbf{A}^{N-2} & \cdots & \mathbf{I} \end{bmatrix} \begin{bmatrix} \mathbf{u}(n) \\ \vdots \\ \mathbf{u}(n+1) \\ \vdots \\ \mathbf{u}(n+N) \end{bmatrix}$$

Extend to Offline Filtering

1. Compute by hopping through $\mathbf{v}(N), \mathbf{v}(2N) \dots$


$$\bullet \quad \mathbf{v}(n+N) = \mathbf{A}^N \mathbf{v}(n) + [\mathbf{A}^{N-1} \quad \mathbf{A}^{N-2} \quad \dots \quad \mathbf{I}] \begin{bmatrix} \mathbf{u}(n) \\ \mathbf{u}(n+1) \\ \vdots \\ \mathbf{u}(n+N-1) \end{bmatrix}$$

2. Compute the rest in parallel:

$$\bullet \quad \begin{bmatrix} \mathbf{v}(n+1) \\ \mathbf{v}(n+2) \\ \vdots \\ \mathbf{v}(n+N-1) \end{bmatrix} = \begin{bmatrix} \mathbf{A} & \mathbf{I} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{A}^2 & \mathbf{A} & \mathbf{I} & \dots & \mathbf{0} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \mathbf{A}^{N-1} & \mathbf{A}^{N-2} & \mathbf{A}^{N-3} & \dots & \mathbf{I} \end{bmatrix} \begin{bmatrix} \mathbf{v}(n) \\ \mathbf{u}(n) \\ \mathbf{u}(n+1) \\ \vdots \\ \mathbf{u}(n+N-2) \end{bmatrix}$$

Note: We can apply step 1 repeatedly (it's a recursive definition).

Assemble Parallel Associative Scan



Step 1 (4 times, N=2)

$O(n) \rightarrow O(\log(n))$

Step 2 (4 times)


Image source: https://en.wikipedia.org/wiki/Prefix_sum

Take away


1. > 1000 small kernels → write custom extension
2. Associative? → parallelisable!
3. You can do both

Benchmarks (M=2, forward)

Intel i7 7700K, 1 thread




RTX 5060 Ti




Benchmarks (M=2, backpropagation)

Intel i7 7700K, 1 thread



RTX 5060 Ti



Tips For Using PhilTorch

- PhilTorch has fast kernel for $M = 1$ or 2 .
 - Sample unrolling for $M > 2$.
- Decompose your filter into second-order sections (SOS)
 - Cascaded SOS 
 - Parallel SOS  
- Check out the example!



 Estimate
lowpass
filter

