

Differentiable Filters: Applications, Advancements, and Challenges

8.1.2025, IIS Academia Sinica

Chin-Yun Yu, Queen Mary University of London

Something about me

- Yu, Chin-Yun (尤靖允)
- B.S. in CS @NYCU (2018)
- RA & Intern @MCTLab (2017~2018)
- Software engineer (day) &
Independent researcher (night)
(2019~2022)
- PhD candidate in AI & Music
@QMUL (2022~)
- Research interests
 - Music Information Retrieval
 - Voice synthesis
 - Audio Effects Modelling
 - Binaural Audio
- Publications
 - ICASSP, ISMIR, DAFX, Interspeech,
WASPAA, AES
- <https://iamycy.github.io/>

What's on the menu?

I will cover:

- Automatic differentiation (AD) in digital filters
- LTI & LPV systems
- GPU acceleration
- PyTorch
- PhilTorch
 - github.com/yoyolicoris/philtorch

I won't cover:

- ❑ Non-linear systems
 - ❑ only a few slides
- ❑ Continuous systems
 - ❑ $H(z)$ 😍
 - ❑ $H(s)$ 😰
- ❑ Gradients of analog filters
 - ❑ No, but that sounds interesting!
- ❑ Applications outside audio
 - ❑ should be generalisable

Outline

- Part I: Digital filters and their differentiable applications
- Part II: Achieving fast automatic differentiation of filters
- Part III: Unsolved challenges and promising directions

Part I: Digital filters and their differentiable applications

What are digital filters?

- In the Z domain

$$H(z) = \frac{b_0 + b_1 z^{-1} + \cdots + b_M z^{-M}}{1 + a_1 z^{-1} + \cdots + a_M z^{-M}}$$

- In the time domain

$$\begin{aligned} y(n) &= b_0 x(n) + b_1 x(n-1) + \cdots + b_M x(n-M) \\ &\quad - a_1 y(n-1) - \cdots - a_M y(n-M) \end{aligned} \quad \xleftrightarrow{\mathcal{Z}} Y(z) = X(z)H(z)$$

Advantages over neural nets

- ★ Readable and controllable
- ★ Low degree of freedom → Embed strong prior knowledge
 - Faster convergence
 - Cheap computation
 - Compact representation
- ★ Widely used in audio applications
 - easier to deploy

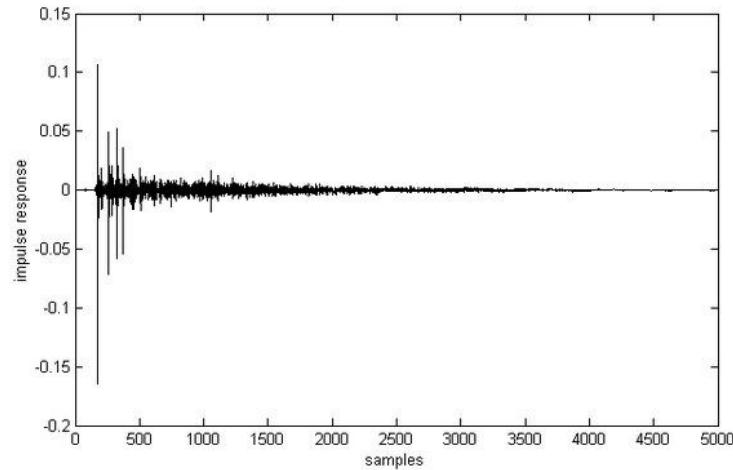
Note: *Still empirical compare to white-box method (e.g., physical modelling)*

Case 1: FIR filters

- Feedforward coefficients only
- Has finite-length impulse response (FIR)

$$\begin{aligned}y(n) &= b_0 x(n) + b_1 x(n - 1) + \dots \\&\quad + b_M x(n - M) \\&= b * x(n)\end{aligned}$$

- ★ Room impulse response (RIR) measurements



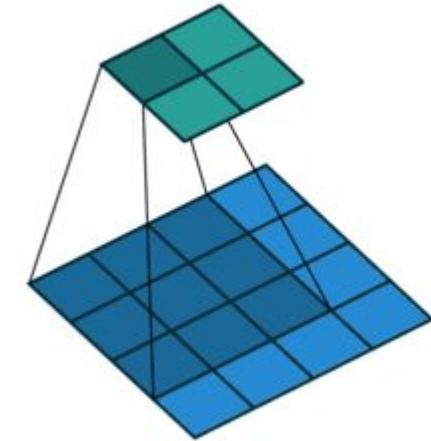
In the deep learning era...

Pros

- Easy to parallelise → Fast on the GPU
- Convolutional neural network (CNN)

Cons

- Limited context window → **limited expressiveness**



End-to-end optimisation

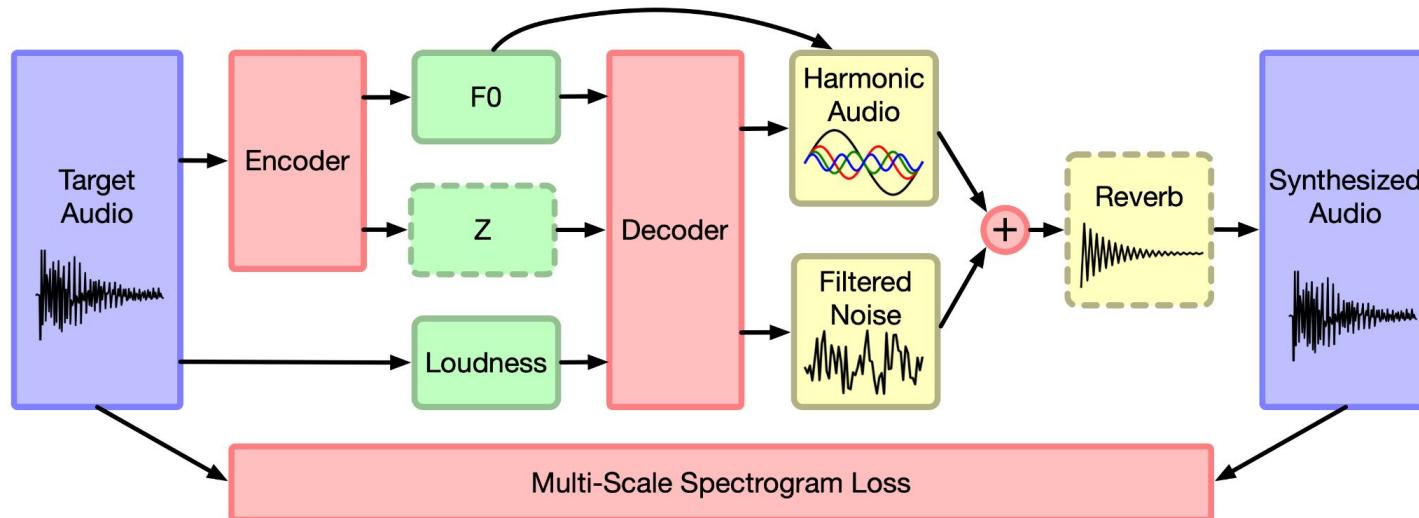
$$\begin{aligned} v(n) &= f_\theta \circ u(n) \\ &= f_1 \circ f_2 \circ \cdots \circ \textcolor{blue}{h} * f_k \circ \cdots \circ u(n) \end{aligned}$$

$$\min_{\theta} \mathcal{L}(v(0), v(1), \dots, v(N-1))$$

$$\theta \mapsto \theta - \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

$$\frac{\partial \mathcal{L}}{\partial \theta} = \cdots + \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial h} \frac{\partial h}{\partial \theta} + \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial \theta} + \cdots$$

Differentiable digital signal processing (DDSP)



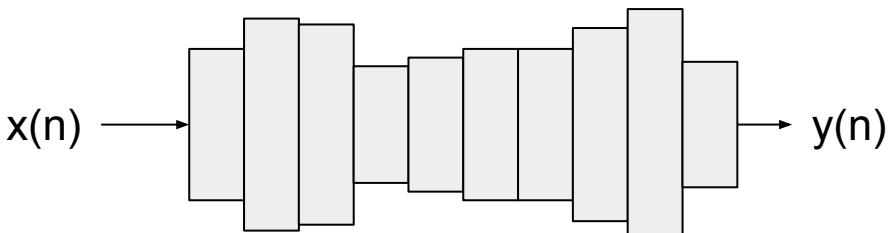
Engel, Jesse, Chenjie Gu, and Adam Roberts. "DDSP: Differentiable Digital Signal Processing." International Conference on Learning Representations.

Case 2: All-pole filters

- a.k.a linear prediction filter
- Feedback coefficients only
- Has infinite-length impulse response (IIR)

- ★ Voice synthesis
 - linear predictive coding (LPC)
- ★ ≈ air flow through a tube with varying diameter
- ★ ≈ resonators

$$\begin{aligned}y(n) &= x(n) - \sum_{i=1}^M a_i y(n-i) \\&= x(n) + \sum_{i=1}^{\infty} h(i)x(n-i)\end{aligned}$$

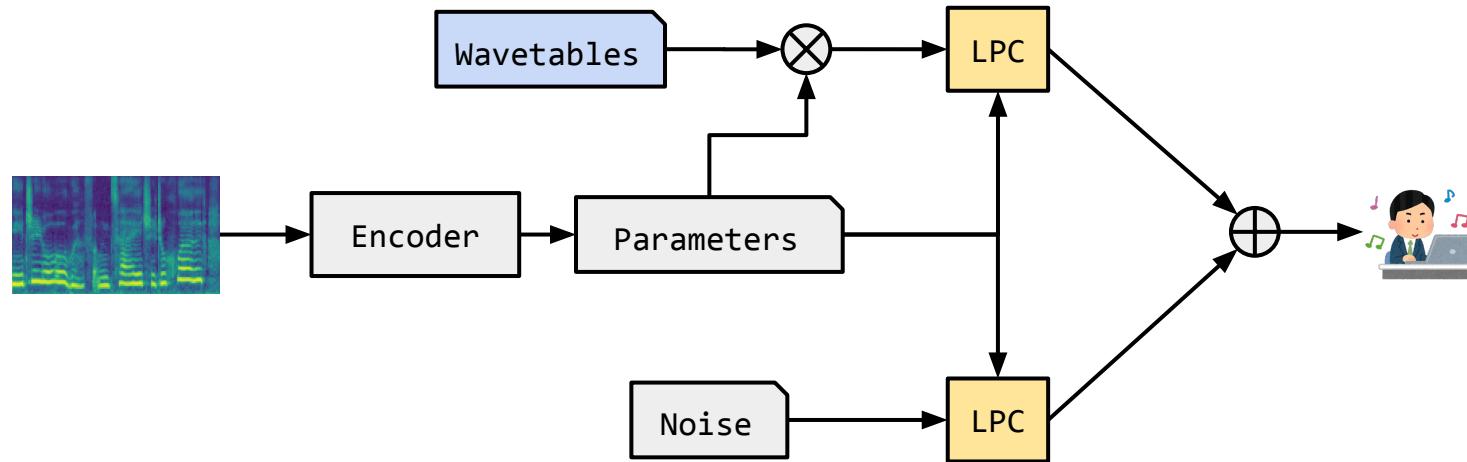




Pink Trombone by Neil Thapen.
<https://dood.al/pinktrombone/>

GIittal-flow LPC Filter (GOLF) vocoder

Yu, Chin-Yun and George Fazekas. "GOLF: A Singing Voice Synthesiser with Glottal Flow Wavetables and LPC Filters." TISMIR.



LP Coefficients Parameterisation

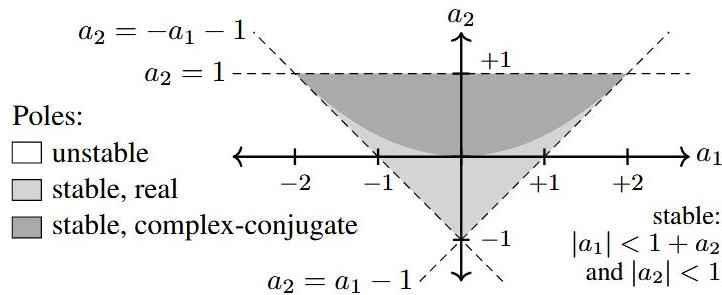


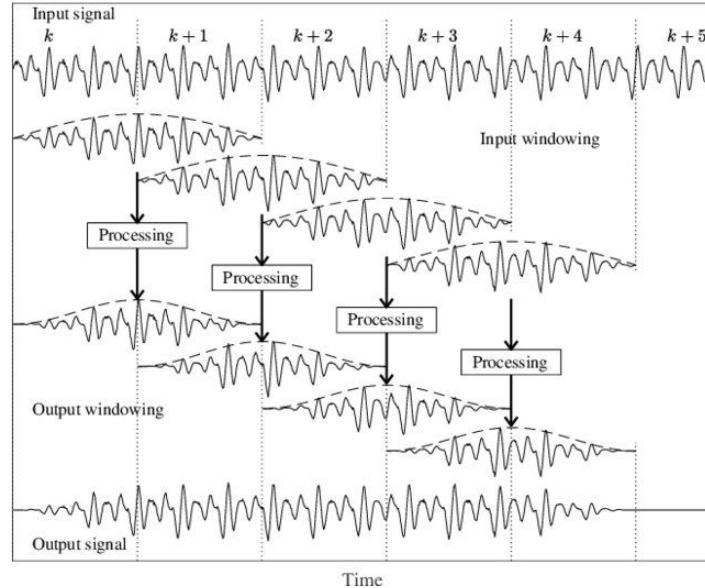
Fig. 1. Biquad triangle ensuring filter stability.

Nercessian, Shahan, Andy Sarroff, and Kurt James Werner.
"Lightweight and interpretable neural modeling of an audio distortion effect using hyperconditioned differentiable biquads." ICASSP 2021.

$$\frac{G}{1+\sum_{i=1}^M a_i z^{-i}} = G \prod_{k=1}^{\frac{M}{2}} \frac{1}{(1+\eta_{k,1}z^{-1}+\eta_{k,2}z^{-2})}$$

- Ensure poles are inside unit circle (stable LTI system)
- $M = 22$, $\text{sr} = 24 \text{ kHz}$

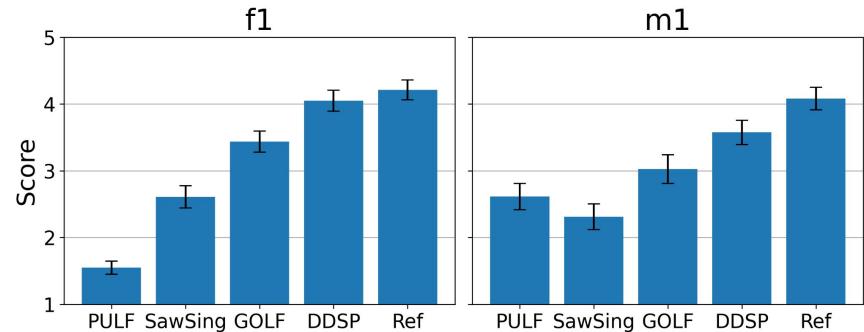
Frame-wise time-varying linear prediction



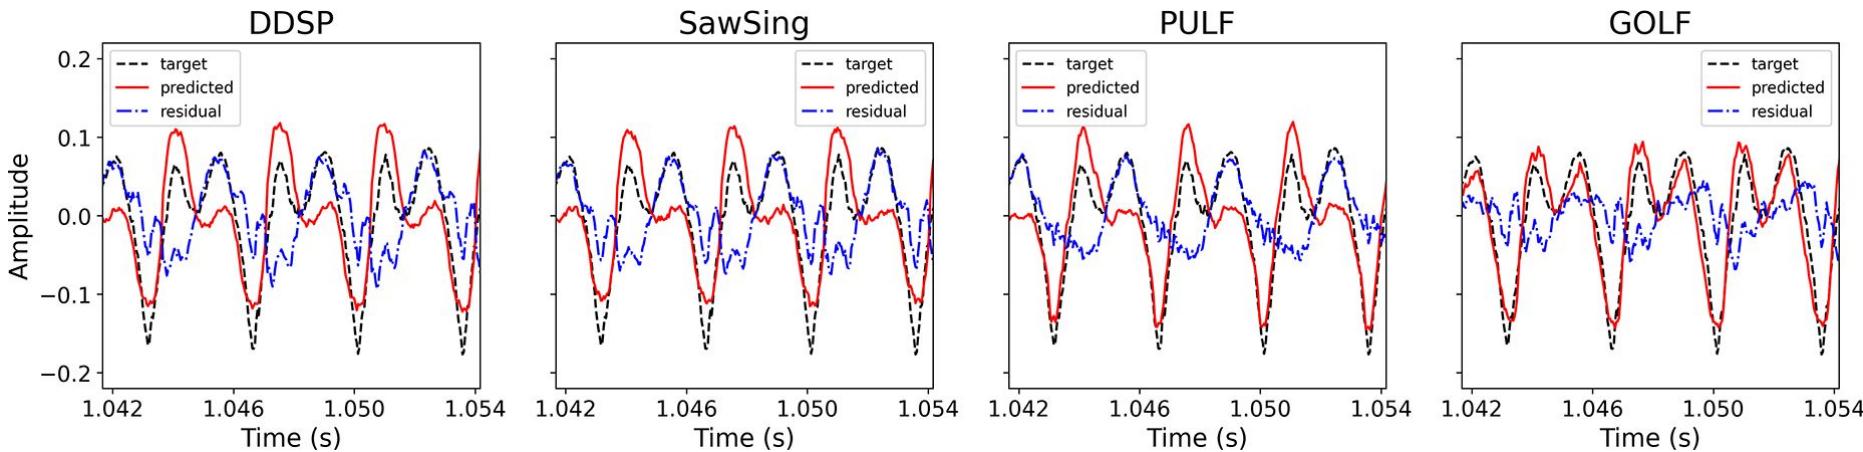
- Overlap-add and windowing
- Each frame is treated as LTI and synthesise **in parallel**

Evaluations

Singers	Models	MSSTFT	MAE-f0 (cent)	FAD
f1	DDSP	3.09	74.47±1.19	0.50 ± 0.02
	SawSing	3.12	78.91 ± 1.18	0.38±0.02
	GOLF	3.21	77.06 ± 0.88	0.62 ± 0.02
	PULF	3.27	76.90 ± 1.11	0.75 ± 0.04
m1	DDSP	3.12	52.95±1.03	0.57 ± 0.02
	SawSing	3.13	56.46 ± 1.04	0.48±0.02
	GOLF	3.26	54.09 ± 0.30	0.67 ± 0.01
	PULF	3.35	54.60 ± 0.73	1.11 ± 0.04



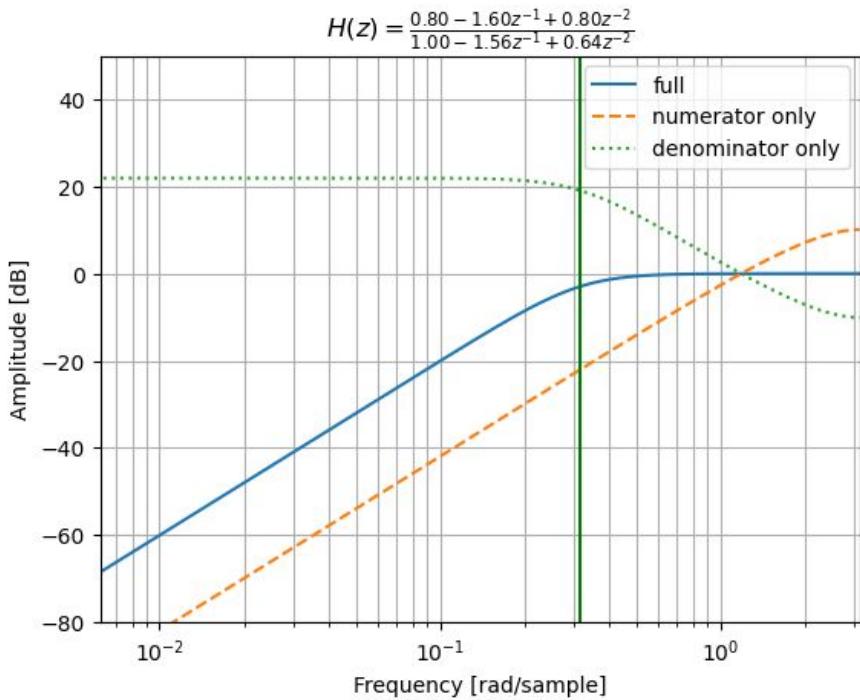
Waveform matching



Case 3: Biquad

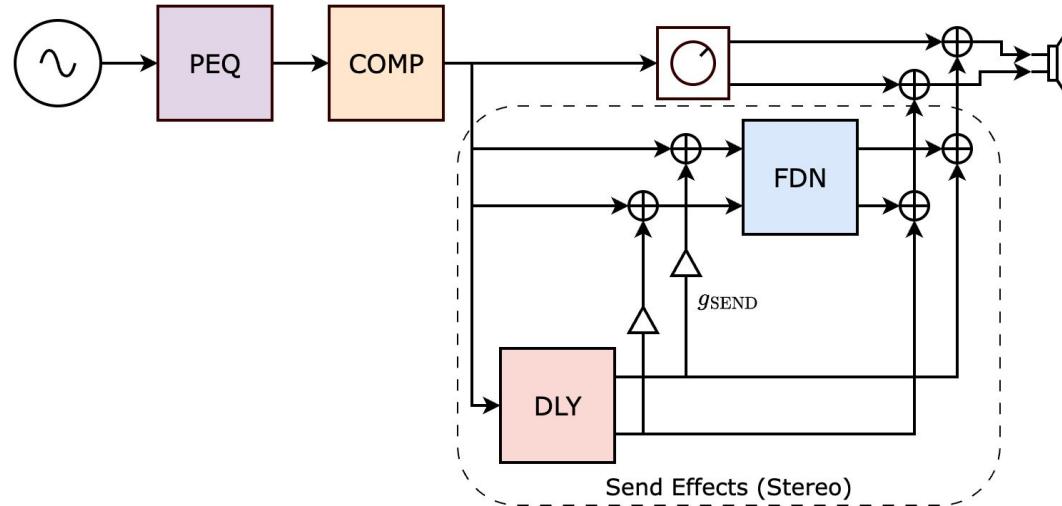
- Mixed feedforward & feedback sections
- Various filters
 - low/high/band pass
 - low/high shelves
 - peak
 - .etc

$$H_{\text{BQ}}(z) = \frac{b_0 + b_1 z^{-1} + b_2 z^{-2}}{1 + a_1 z^{-1} + a_2 z^{-2}}$$



Differentiable Vocal Fx (DiffVox)

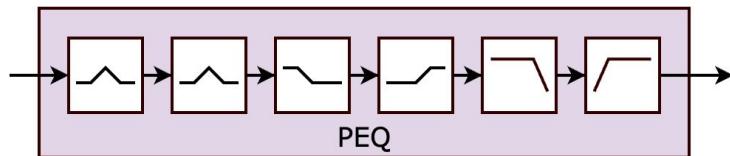
Yu, Chin-Yun, et al. "DiffVox: A differentiable model for capturing and analysing vocal effects distributions." Proc. DAFX. 2025.



Sony AI

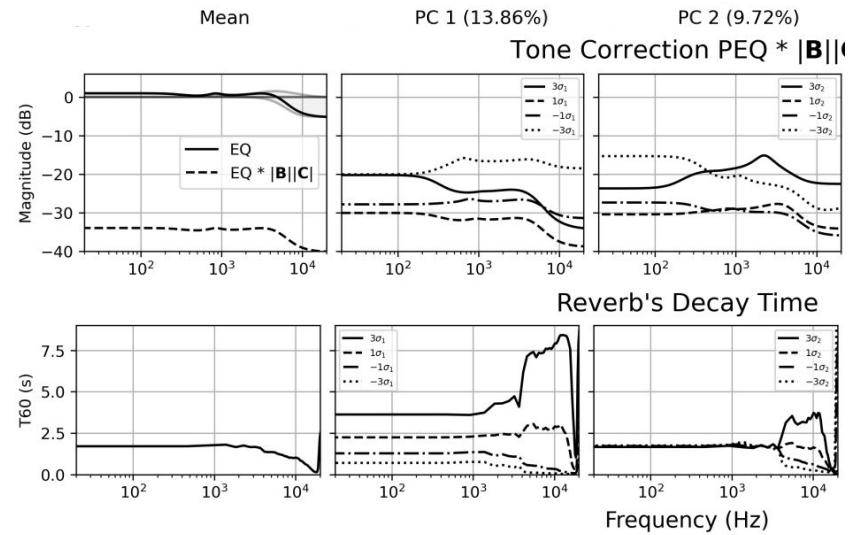
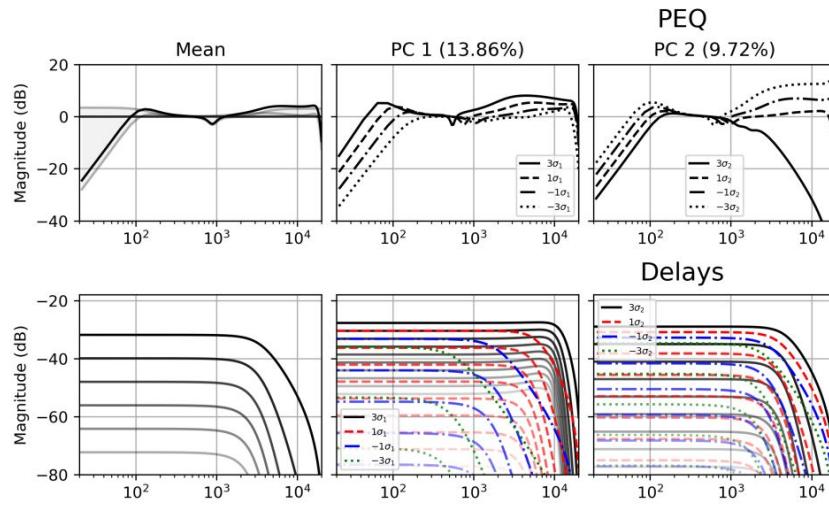
6-band parametric equaliser (PEQ)

- Biquad filters
- peak filter x2, low-shelf, high-shelf, low-pass, high-pass
- ≈ IK Multimedia Classic Equalizer VST



Principal component analysis (PCA)

Fit DiffVox on Sony's 365 proprietary vocal tracks



Recursive filters (IIR)

Pros

- Unlimited context window → extremely expressive with a few coefficients
- Physical meaning (e.g., formant frequencies)
- Friendly to audio engineers

Cons

- Sequential computation (?) → non-trivial to parallelise (**are they?**)
- Can be unstable and care should be taken

Part II: Achieving fast automatic differentiation of filters

Issue in the most popular AD framework

- Popularity: PyTorch >>> JAX ≈ Tensorflow > etc.
- No naive ops for recursion → Slow to compute
 - ML people understand it as autoregression

$$y(n) = b_0x(n) + b_1x(n - 1) + \cdots + b_Mx(n - M)$$
$$- a_1y(n - 1) - \cdots - a_My(n - M)$$

`torch.nn.functional.conv1d`
`torch.* ?`

A naive implementation

```
h = zi
for xn in x.unbind(1):
    h = torch.addmm(xn, h, AT)
    results.append(h if out_idx is None else h[:, out_idx])
output = torch.stack(results, dim=1)
```

philtorch/lti/ssm.py, L127-131

runtime



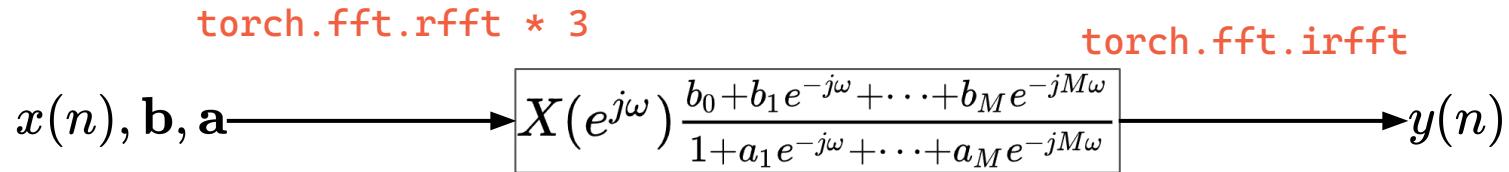
F.conv1d

kernel

Just-in-time (JIT) compilers

- Freeze the computation graph to optimise further
- `torch.jit.script/trace`
 - 2 ~ 3 times faster
 - Deprecated $\geq v2.9$
- `torch.compile`
 - for-loop? **unroll them!** 😊
 - million audio samples → your programme hang 4ever **during compilation**
 - Wu, Yidi, et al. "Control Flow Operators in PyTorch." 2025.
- Fundamental cause (personal opinion)
 - Eager mode first & entry-level friendly to beginner

Frequency sampling (FS)

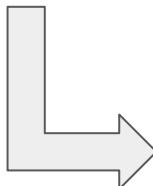


```
y = irfft(  
    rfft(x) * rfft(b, n=x.shape[0]) / rfft(a, n=x.shape[0]),  
    n=x.shape[0]  
)
```

Time aliasing

- S&S 101: **Sampling → periodic response**
- Duplicates IIR at location ..., -2N, -N, 0, N, 2N, ...
 - = folded and wrapped inside $n \in [0, N]$

$$y(n) = x(n) + \alpha y(n-1) \rightarrow \mathbf{h} = [1, \alpha, \alpha^2, \dots]^{\top}$$

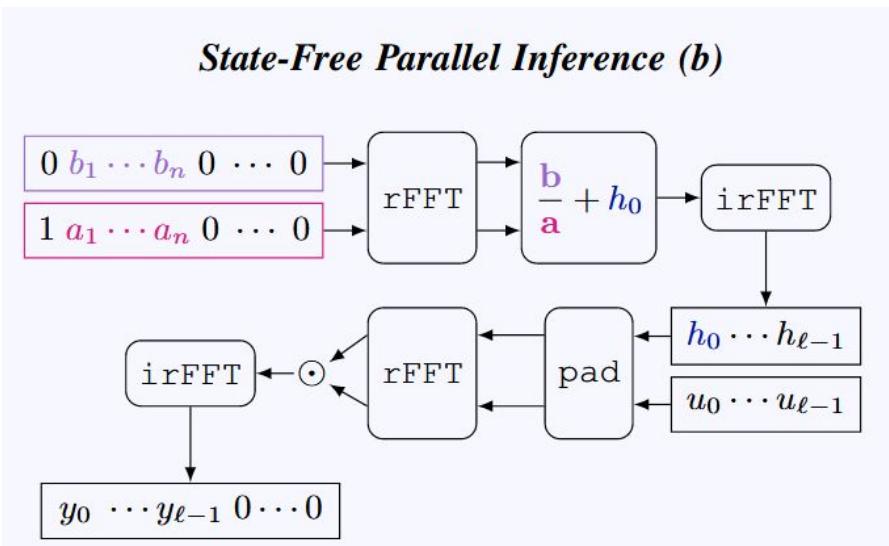
Frequency sampling 

$$\begin{aligned}\hat{\mathbf{h}} = & [1, \alpha, \alpha^2, \dots, \alpha^{N-1}]^{\top} \\ & + [\alpha^N, \alpha^{N+1}, \alpha^{N+2}, \dots, \alpha^{2N-1}]^{\top} \\ & + [\alpha^{2N}, \alpha^{2N+1}, \alpha^{2N+2}, \dots, \alpha^{3N-1}]^{\top} \\ & + \dots\end{aligned}$$

Circular convolution

- Periodic IR → circular convolution
- FIR: zero-padding to $2N - 1$
- IIR: how? there's no end point!

Solution: IIR truncation



$$H(z) = \mathbf{C}(z\mathbf{I} - \mathbf{A})^{-1}\mathbf{B} + b_0$$
$$H_N(z) = \mathbf{C}(\mathbf{I} - \mathbf{A}^N)(z\mathbf{I} - \mathbf{A})^{-1}\mathbf{B} + b_0$$
$$\xleftarrow{z} \mathbf{h}_N = [1, \alpha, \dots, \alpha^{N-1}]^\top$$
$$\tilde{\mathbf{b}} = \mathbf{C}(\mathbf{I} - \mathbf{A}^N)$$

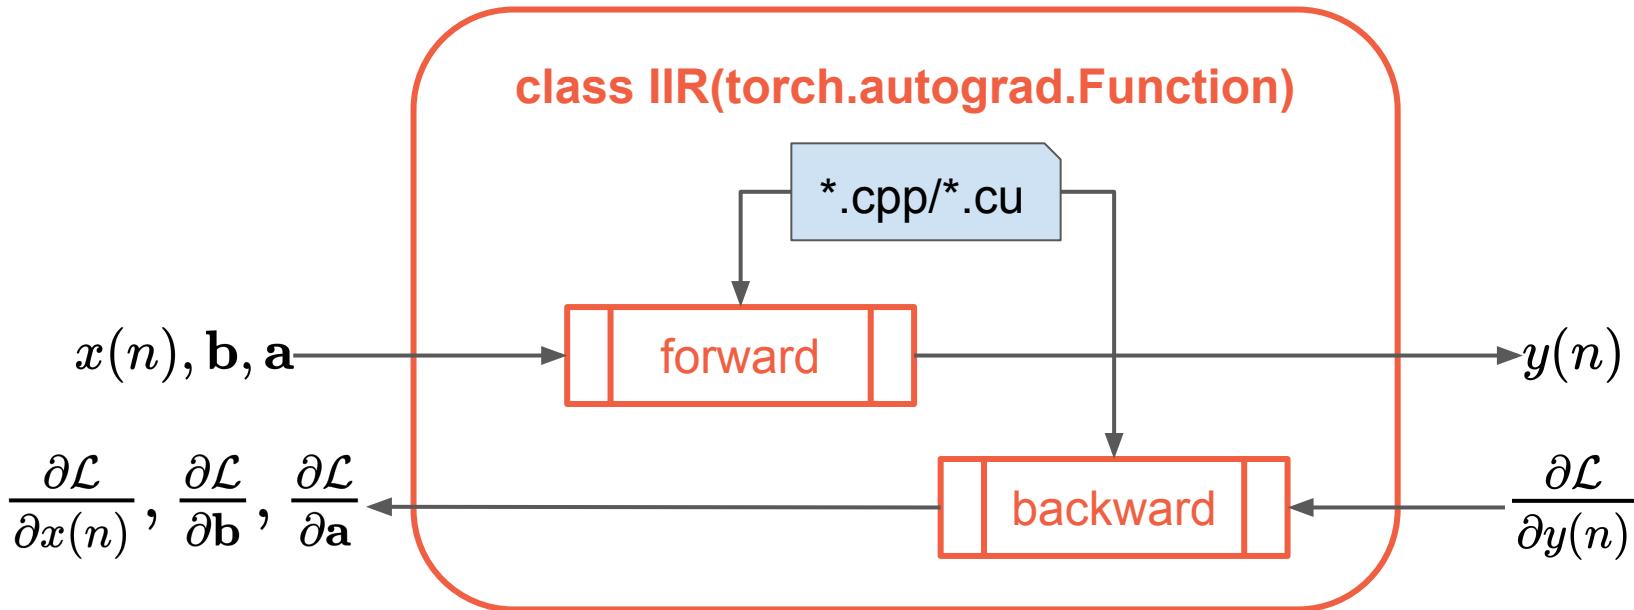
- **A:** companion matrix of the **a** polynomial
- Evaluate matrix exponential \mathbf{A}^N

Parnichkun, Rom N., et al. "State-free inference of state-space models: the transfer function approach." ICML. 2024.

Time domain filtering

- FS
 - Overkill when $M \ll N$
 - Sensitive to numerical errors
 - Higher memory usage
- Direct computation is favourable for most filters

Make filter the first class citizen



https://intro2ddsp.github.io/filters/iir_torch.html

Step 1: Analytical gradients

Gradient backpropagation

Reverse-mode AD

$$\frac{\partial \mathcal{L}}{\partial \theta} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial b} \frac{\partial b}{\partial \theta} + \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial a} \frac{\partial a}{\partial \theta} + \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial \theta} + \dots$$

- Full jacobian : expensive
- Vector-jacobian product 

$$\text{VJP}_{\mathbf{J}} : \mathbf{v} \mapsto \mathbf{v}\mathbf{J}$$

$$\frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial b} \frac{\partial b}{\partial \theta} = \left(\text{VJP}_{\frac{\partial b}{\partial \theta}} \circ \text{VJP}_{\frac{\partial y}{\partial b}} \circ \text{VJP}_{\frac{\partial \mathcal{L}}{\partial y}} \right) (1)$$

- Compute order: →
- Apply transformation **without transformation matrices**

Filtering as matrix transformation

$$\mathbf{y} = \mathbf{W}\mathbf{x}$$

$$\mathbf{W} = \begin{bmatrix} h(0) & 0 & 0 & \cdots & 0 \\ h(1) & h(0) & 0 & \cdots & 0 \\ h(2) & h(1) & h(0) & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ h(-N+1) & h(-N+2) & h(-N+3) & \cdots & h(0) \end{bmatrix}$$

Matrix calculus 101

$$\mathbf{y} = \mathbf{W}\mathbf{x}$$

↓

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \mathbf{W} = \left(\mathbf{W}^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}^\top} \right)^\top$$
$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{W}} = \mathbf{x} \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$$

forward filtering

reverse filtering

The two-pass algorithm: for \mathbf{x}

$$\mathbf{W}\mathbf{x} \longleftrightarrow \text{AllPole}(\mathbf{x})$$

$$\mathbf{W}^\top \frac{\partial \mathcal{L}}{\partial \mathbf{y}^\top} \longleftrightarrow (\text{Flip} \circ \text{AllPole} \circ \text{Flip}) \left(\frac{\partial \mathcal{L}}{\partial \mathbf{y}^\top} \right)$$

Yu, Chin-Yun, and György Fazekas. "Singing Voice Synthesis Using Differentiable LPC and Glottal-Flow-Inspired Wavetables." ISMIR. 2023.

Forgione, Marco, and Dario Piga. "dynoNet: A neural network architecture for learning dynamical systems." International Journal of Adaptive Control and Signal Processing 35.4 (2021): 612-626.

The two-pass algorithm: for a

$$\frac{\partial y(n)}{\partial a_i} = -y(n-i) - \sum_{k=1}^M a_k \frac{\partial y(n-k)}{\partial a_i}$$

$$\frac{\partial \mathbf{y}}{\partial \mathbf{a}} = \text{AllPole}(-\mathbf{y})$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{a}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{a}}$$

- Filtering x2
- Matrix multiplications x1

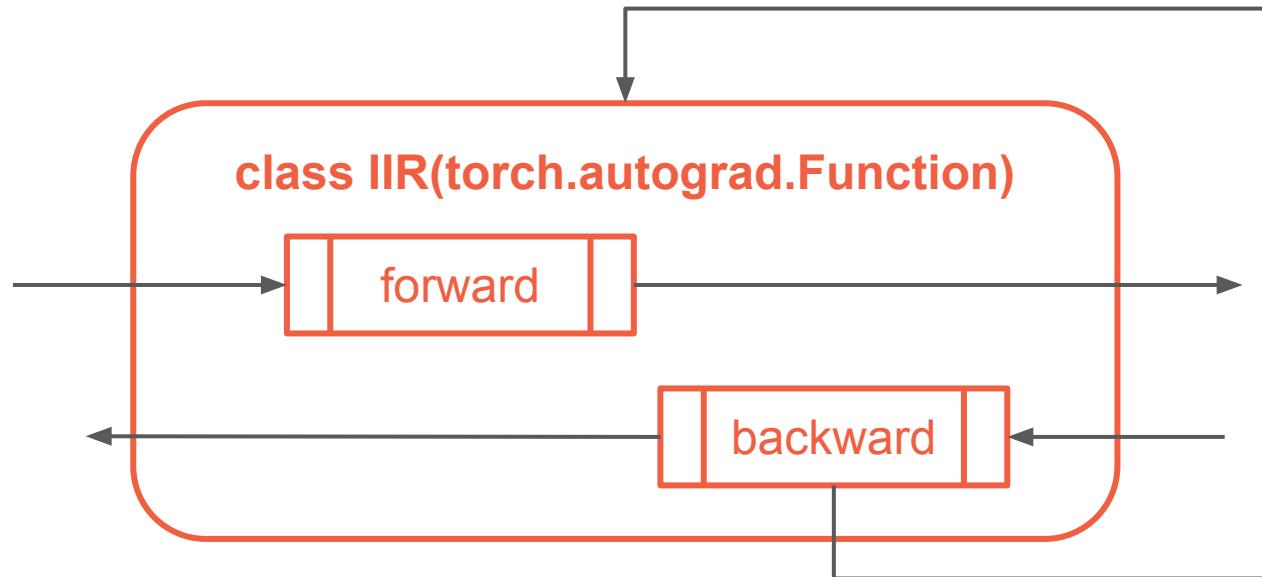
The one-pass algorithm

$$\frac{\partial \mathcal{L}}{\partial \mathbf{a}} = -\frac{\partial \mathcal{L}}{\partial \mathbf{x}} \begin{bmatrix} y(-1) & y(-2) & \cdots & y(-M) \\ y(0) & y(-1) & \cdots & y(-M+1) \\ \vdots & \vdots & \ddots & \vdots \\ y(N-2) & y(N-3) & \cdots & y(N-M-1) \end{bmatrix}$$

- Filtering x1
- Matrix multiplications x1
- TorchAudio `lfilter`

Yu, C.Y., Mitcheltree, C., Carson, A., Bilbao, S., Reiss, J.D. and Fazekas, G., 2024. Differentiable All-pole Filters for Time-varying Audio Systems. DAFX 2024.

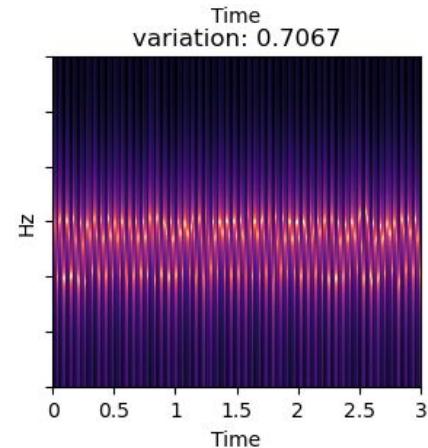
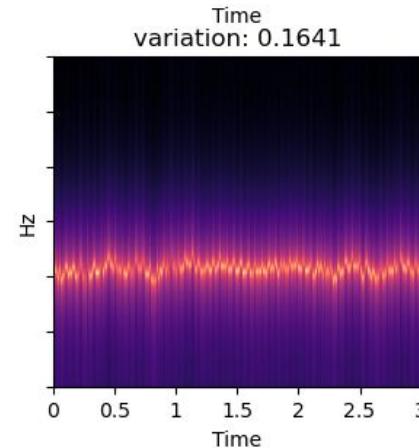
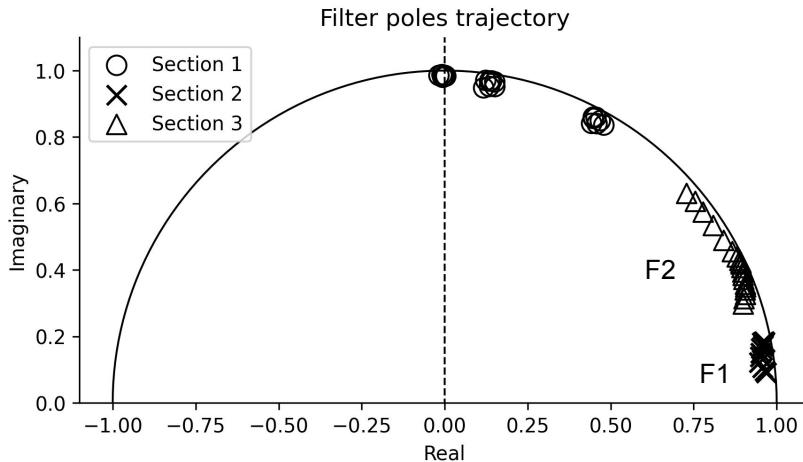
N-order gradients are recursively defined



You only implement once (YOIO)

Extending to parameter-varying systems

~ to remove frame-wise approximation errors in GOLF



- Each point is covered by four frames (75 % overlap)
- Poles are interleaving between three groups (<4)

Yu, Chin-Yun and George Fazekas. "GOLF: A Singing Voice Synthesiser with Glottal Flow Wavetables and LPC Filters." TISMIR.

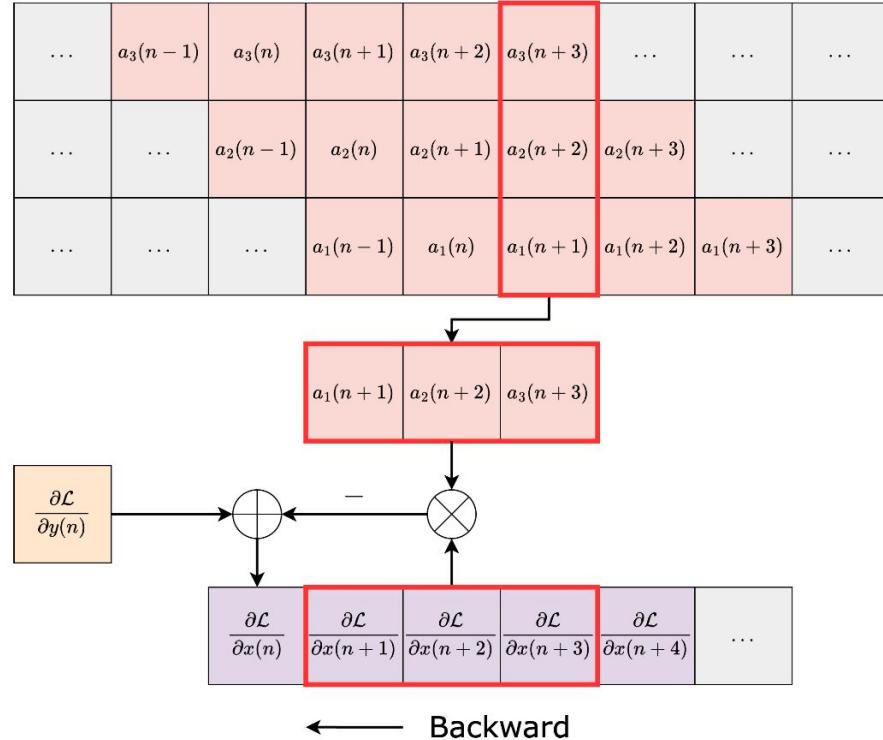
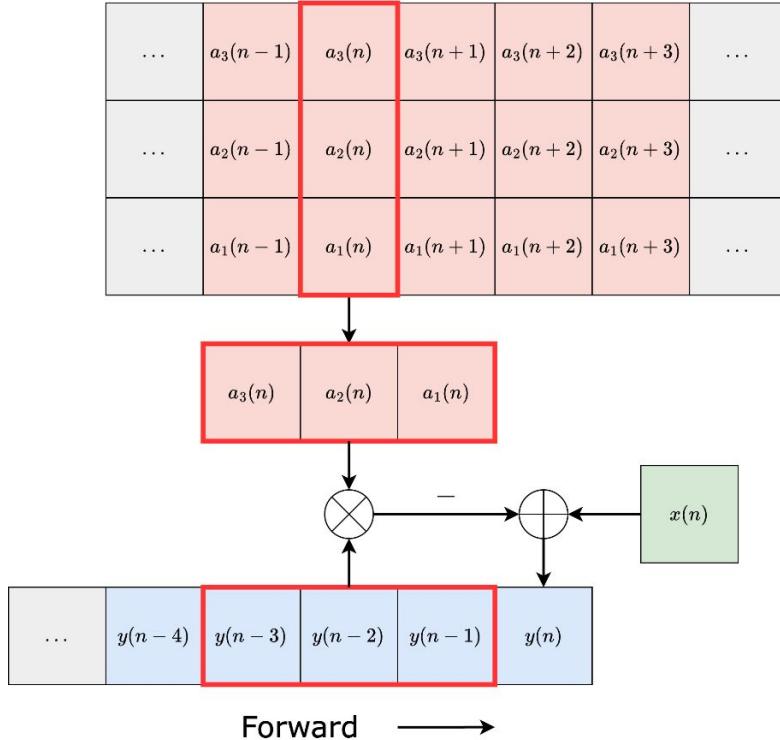
Gradients of a LPV all-pole

$$y(n) = x(n) - \sum_{i=1}^M a_i(n)y(n-i)$$

$$\frac{\partial \mathcal{L}}{\partial x(n)} = \frac{\partial \mathcal{L}}{\partial y(n)} - \sum_{i=1}^M a_i(\textcolor{blue}{n+i}) \frac{\partial \mathcal{L}}{\partial x(n+i)}$$

$$\frac{\partial \mathcal{L}}{\partial a_i(n)} = -\frac{\partial \mathcal{L}}{\partial x(n)} y(n-i)$$

Yu, C.Y., Mitcheltree, C., Carson, A., Bilbao, S., Reiss, J.D. and Fazekas, G., 2024. Differentiable All-pole Filters for Time-varying Audio Systems. DAFX 2024.

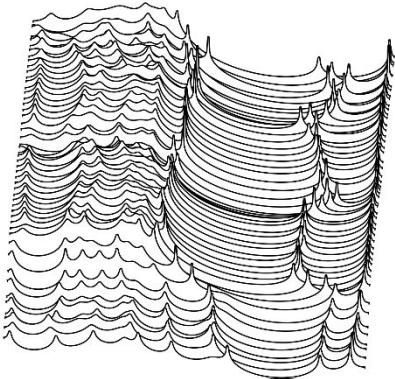


Yu, C.Y., Mitcheltree, C., Carson, A., Bilbao, S., Reiss, J.D. and Fazekas, G., 2024. Differentiable All-pole Filters for Time-varying Audio Systems. DAFX 2024.

A better GOLF vocoder

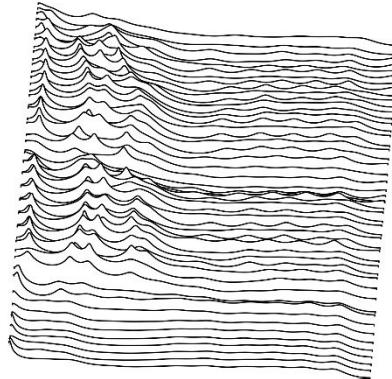


GOLF-v1



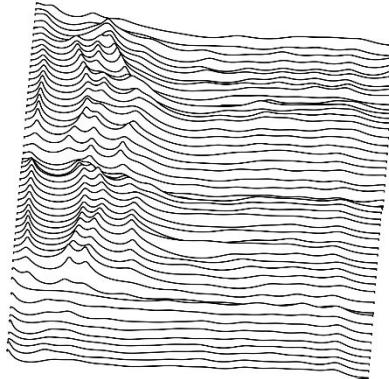
Harmonic-plus-noise
frame-wise LP

GOLF-ff



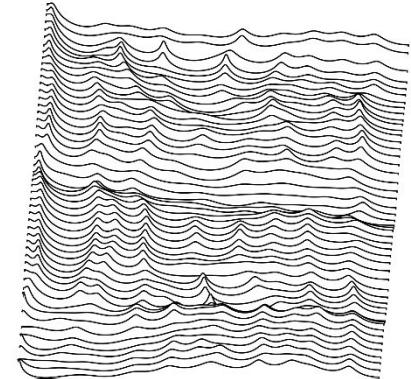
Source-filter
frame-wise LP

GOLF-ss



Source-filter
sample-wise LP

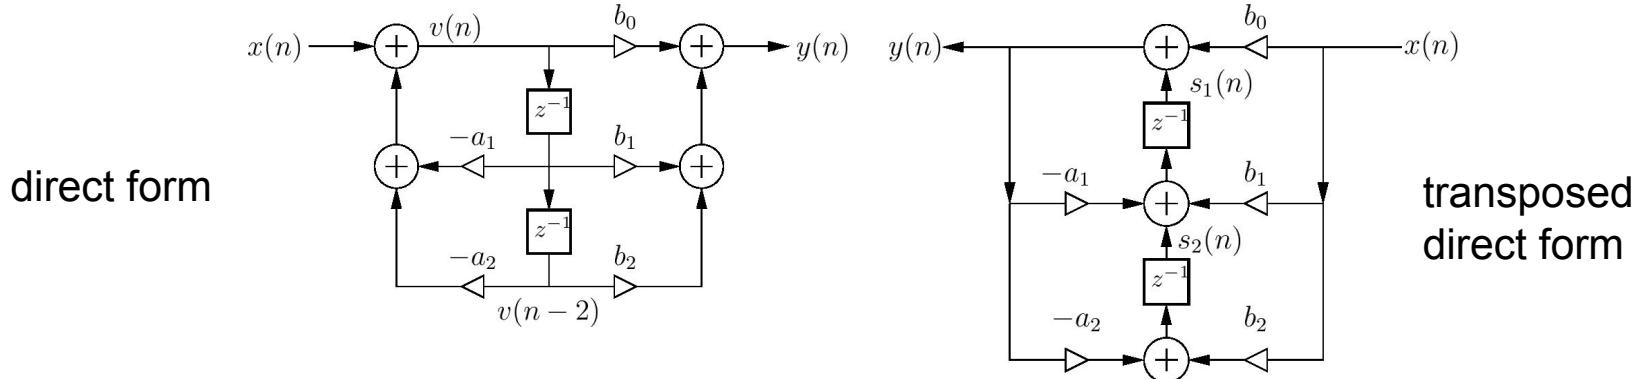
SPTK



Yu, Chin-Yun, and György Fazekas. "Differentiable Time-Varying Linear Prediction in the Context of End-to-End Analysis-by-Synthesis." Proc. Interspeech 2024.

Limitations

- Not all direct forms are FIR/all-pole separable
 - DF-I/II
 - TDF-I/II
- Sequential computation → not fully utilise GPU



The state-space realisation (of DF-II)

$$\mathbf{v}(n+1) = \mathbf{A}\mathbf{v}(n) + \mathbf{B}x(n)$$

$$y(n) = \mathbf{C}^\top \mathbf{v}(n) + b_0 x(n)$$

$$\mathbf{B} = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \mathbf{C} = \begin{bmatrix} b_1 - a_1 b_0 \\ b_2 - a_2 b_0 \\ \vdots \\ b_M - a_M b_0 \end{bmatrix}, \mathbf{A} = \begin{bmatrix} -\mathbf{a} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}$$

TDF-II state-space form

$$\begin{aligned}\mathbf{s}(n+1) &= \mathbf{A}^{\top} \mathbf{s}(n) + \mathbf{C}x(n) \\ y(n) &= \mathbf{B}^{\top} \mathbf{v}(n) + b_0 x(n)\end{aligned}$$

Duality: The two sides of the same coin

TDF-II, incremental in time

$$\mathbf{v}(n+1) = \mathbf{A}\mathbf{v}(n) + \mathbf{u}(n)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n)} = \mathbf{A}^\top \frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n+1)} + \frac{\partial \mathcal{L}}{\partial \mathbf{v}^\top(n+1)}$$

DF-II, decremental in time

DF-II, incremental in time

$$\mathbf{s}(n+1) = \mathbf{A}^\top \mathbf{s}(n) + \mathbf{u}(n)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n)} = \mathbf{A} \frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n+1)} + \frac{\partial \mathcal{L}}{\partial \mathbf{s}^\top(n+1)}$$

TDF-II, decremental in time

Yu, Chin-Yun, and György Fazekas. "Accelerating Automatic Differentiation of Direct Form Digital Filters." DlffSys Workshop @EurIPS 2025.

LPV state-space

$$\mathbf{v}(n+1) = \mathbf{A}(n)\mathbf{v}(n) + \mathbf{u}(n)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n)} = \mathbf{A}^\top(n) \frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n+1)} + \frac{\partial \mathcal{L}}{\partial \mathbf{v}^\top(n+1)}$$

$$\mathbf{A}, a_i(n) \longleftrightarrow \mathbf{A}^\top, a_i(n+i)$$

<https://iamycy.github.io/posts/2025/04/differentiable-tdf-ii/>

Step 2: Parallel implementations

SSM is parallelisable!

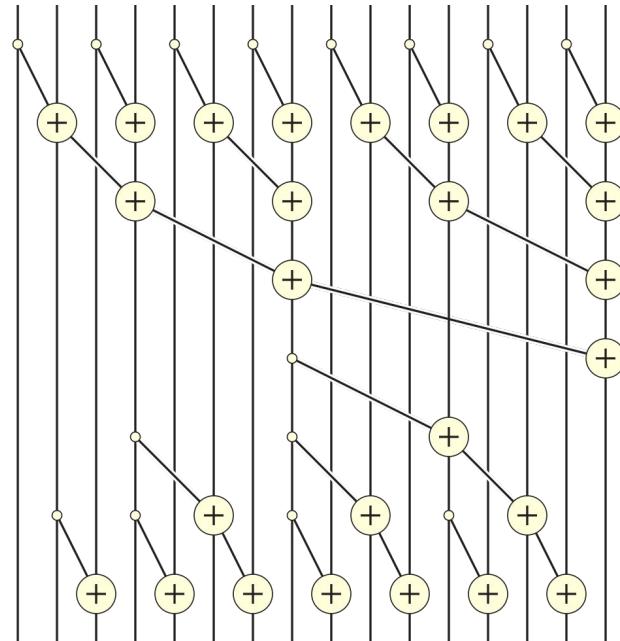
Sequential so non-parallelisable 😞

$$\begin{aligned}\mathbf{v}(n+1) = & \mathbf{A}(n)(\mathbf{A}(n-1)(\mathbf{A}(n-2)(\dots) \\ & + \mathbf{u}(n-2)) + \mathbf{u}(n-1)) + \mathbf{u}(n)\end{aligned}$$

Parallelisable ❤️

$$\begin{aligned}(\mathbf{U}, \mathbf{x}) \oplus (\mathbf{V}, \mathbf{z}) &\mapsto (\mathbf{V}\mathbf{U}, \mathbf{V}\mathbf{x} + \mathbf{z}) \\ (\mathbf{0}, \mathbf{v}(n+1)) = & (\mathbf{0}, \mathbf{v}(0)) \oplus (\mathbf{A}(0), \mathbf{u}(0)) \oplus (\mathbf{A}(1), \mathbf{u}(1)) \\ & \oplus \dots \oplus (\mathbf{A}(N), \mathbf{u}(N))\end{aligned}$$

Parallel associative scan



- Prefix sum/inclusive scan
- `std::inclusive_scan`
- `thrust::inclusive_scan`
- Doable using `torch.matmul`
 - <https://iamycy.github.io/posts/2025/06/28/unroll-ssm/>
- $O(N) \rightarrow O(\log(N))$

https://en.wikipedia.org/wiki/Prefix_sum

Even faster: Diagonalised SSM (LTI only)

$$\tilde{\mathbf{v}}(n+1) = \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_M \end{bmatrix} \tilde{\mathbf{v}}(n) + \mathbf{P}^{-1} \mathbf{u}(n)$$
$$\mathbf{v}(n) = \mathbf{P} \tilde{\mathbf{v}}(n), \quad \mathbf{A} = \mathbf{P} \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_M \end{bmatrix} \mathbf{P}^{-1}$$

Yu, Chin-Yun, et al. "DiffVox: A differentiable model for capturing and analysing vocal effects distributions." Proc. DAFX. 2025.

PFE interpretation of diagonalised SSM

- Partial Fraction Expansion
- M first-order filters run in parallel
 - or M/2 biquads

$$\begin{aligned} H(z) &= b_0 + z^{-1} \sum_{i=1}^M \frac{\tilde{c}_i \tilde{b}_i}{1 - \lambda_i z^{-1}} \\ &= b_0 + \sum_{i=1}^{\frac{M}{2}} \frac{r_{i,1} z^{-1} + r_{i,2} z^{-2}}{1 + a_{i,1} z^{-1} + a_{i,2} z^{-2}} \end{aligned}$$

Part II Summary

- Number of kernel calls is **input-dependent** → pack them into one kernel
 - so it is **input-independent**
- Backpropagation of filters = filter the gradients backward in time
- Linear **recursive filters are parallelisable** across time dimension

PhilTorch Φ 🔥

- Fast and differentiable digital filters/SSMs on GPU
- Same interface and numerical robustness as SciPy
- LTI & LPV systems
- Arbitrary orders of gradients

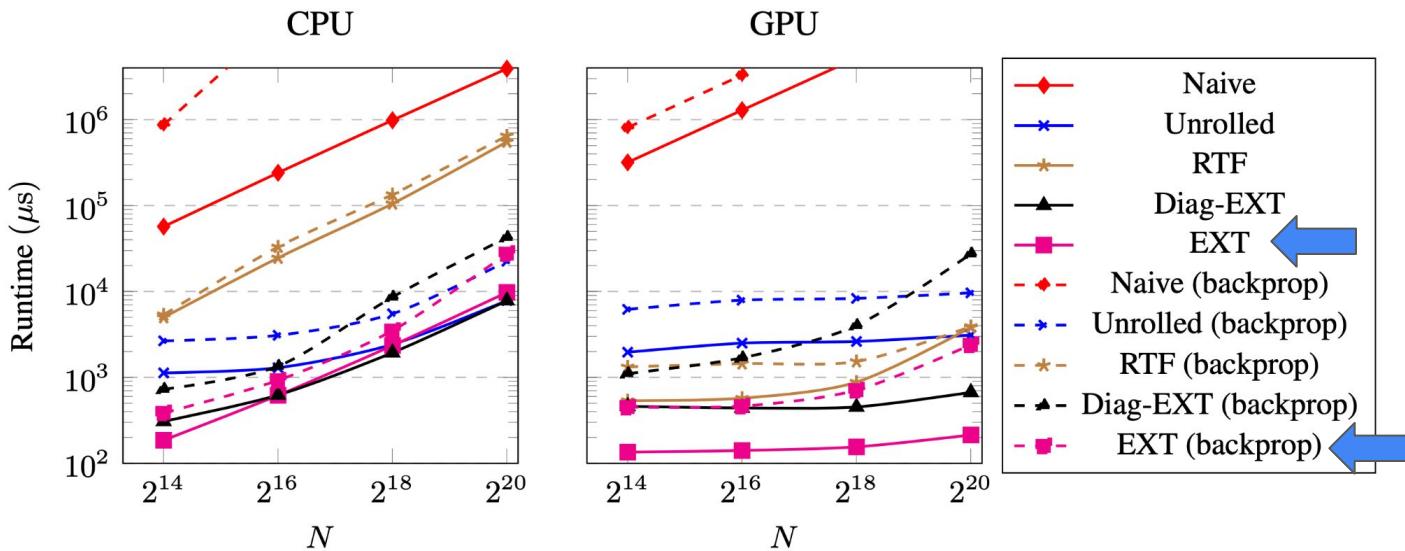
```
pip install philtorch
```



Code



Benchmarks (biquad)



Yu, Chin-Yun, and György Fazekas. "Accelerating Automatic Differentiation of Direct Form Digital Filters." DiffSys Workshop @EurIPS 2025.

Part III: Unsolved challenges and promising directions

Topics

1. Tasks that have parallel filters
2. Non-linear filters
3. Forward-mode gradient computations
4. Learnable delay lengths
5. FDN
6. LPV stability
7. Gradients' numerical accuracy

Parallel form >>> others

- For fast development we need **Speed!!!**
- Work on tasks that have parallel forms
- Avoid repeated poles as much as possible



$$H(z) = b_0 + \sum_{i=1}^{\frac{M}{2}} \frac{r_{i,1}z^{-1} + r_{i,2}z^{-2}}{1 + a_{i,1}z^{-1} + a_{i,2}z^{-2}}$$

Example 1: Feedback comb filter

$$y(n) = x(n) - \alpha y(n - M)$$

↓

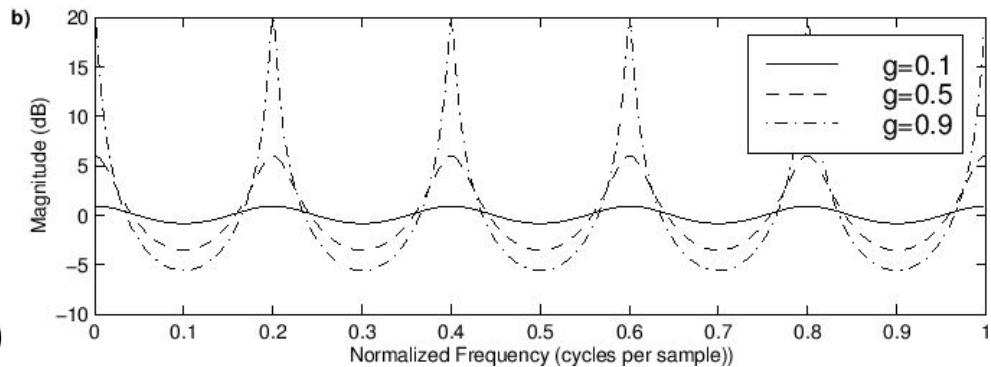
$$y_1(n) = x(nM) - \alpha y_1(n - 1)$$

$$y_2(n) = x(nM + 1) - \alpha y_2(n - 1)$$

⋮

$$y_M(n) = x(nM + M - 1) - \alpha y_M(n - 1)$$

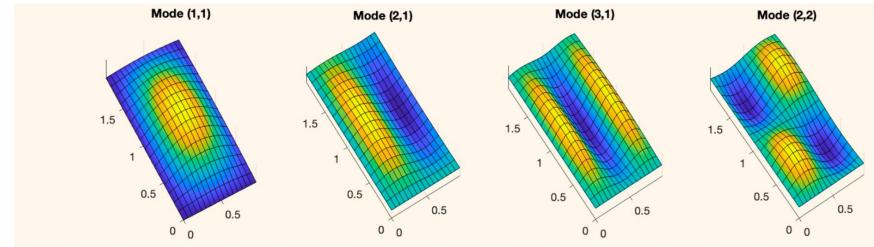
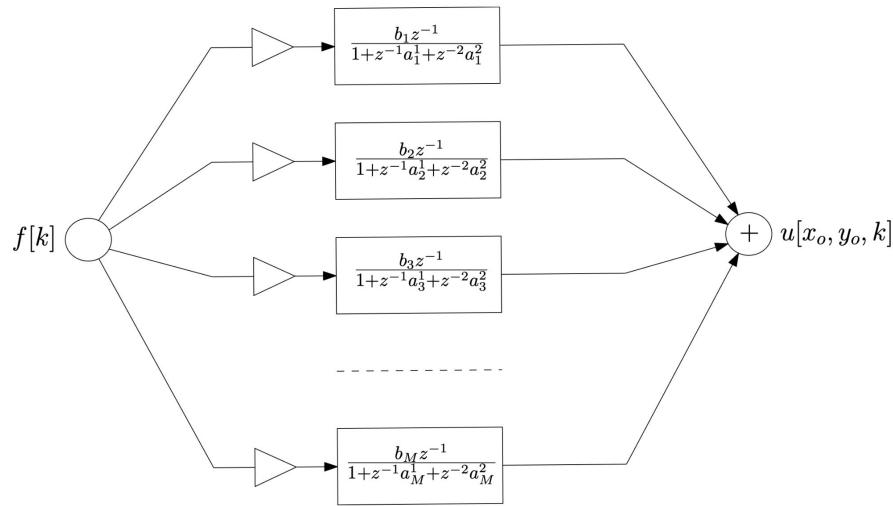
$\mathbf{y} = \text{Interleave}(\mathbf{y}_1, \dots, \mathbf{y}_M)$



`philtorch.lti.comb_filter`

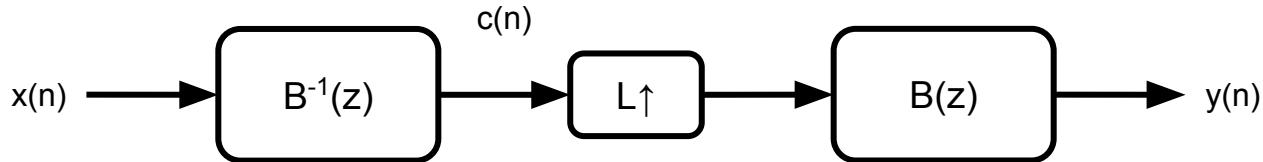
https://www.dsprelated.com/freebooks/pasp/Comb_Filters.html

Example 2: Modal synthesis of a plate reverb



- 1st DAFX Parameter Estimation Challenge
 - now ~ 31.5.2026
- <https://github.com/LOGUNIVPM/1st-DAFx-Challenge>

Example 3: Cubic spline interpolation



$$\begin{aligned}B^{-1}(z) &= \frac{6}{z + 4 + z^{-1}} \\&= \frac{-6\alpha}{(1 - \alpha z^{-1})(1 - \alpha z)} \\&= \frac{-6\alpha}{1 - \alpha^2} \left(\frac{1}{1 - \alpha z^{-1}} + \frac{1}{1 - \alpha z} - 1 \right)\end{aligned}$$

`philtorch.lti.cubic_spline`

Unser, Michael, Akram Aldroubi, and Murray Eden. "Fast B-spline transforms for continuous image representation and interpolation." IEEE Transactions on pattern analysis and machine intelligence 13.3 (1991): 27

Accelerate AD of non-linear filters

- Their forward pass cannot be parallelised
- ...but their backprop can!
- Challenges: dense Jacobians

$$\begin{aligned}\mathbf{v}(n+1) &= f(\mathbf{v}(n)) + \mathbf{u}(n) \\ \frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n)} &= \sum_{i=n+1}^N \left(\prod_{j=n+1}^{i-1} \left(\frac{\partial \mathbf{v}(j+1)}{\partial \mathbf{v}(j)} \right)^\top \right) \frac{\partial \mathcal{L}}{\partial \mathbf{v}^\top(i)} \\ &= \left(\frac{\partial \mathbf{v}(n+2)}{\partial \mathbf{v}(n+1)} \right)^\top \frac{\partial \mathcal{L}}{\partial \mathbf{u}^\top(n+1)} + \frac{\partial \mathcal{L}}{\partial \mathbf{v}^\top(n+1)}\end{aligned}$$

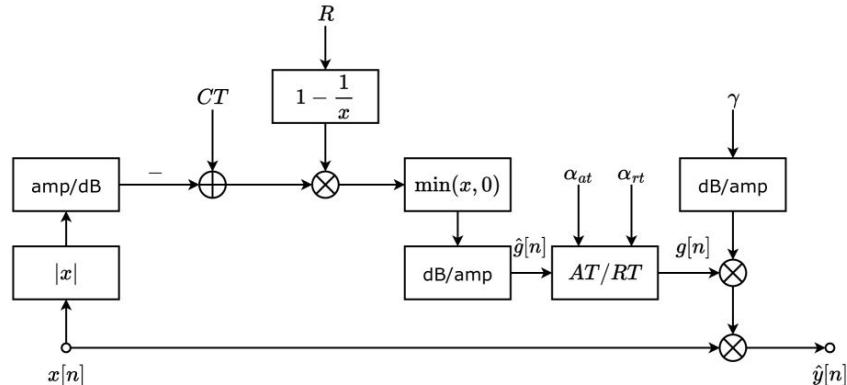
https://justintchiu.com/blog/pscan_diff/

Wang, Shang, Yifan Bai, and Gennady Pekhimenko. "BPPSA: Scaling back-propagation by parallel scan algorithm." Proceedings of Machine Learning and Systems 2 (2020): 451-469.

Example: Differentiable gain reduction in a compressor

$$g(n) = \begin{cases} \alpha_{at}\hat{g}(n) + (1 - \alpha_{at})g(n - 1) & \text{if } \hat{g}(n) < g(n - 1), \\ \alpha_{rt}\hat{g}(n) + (1 - \alpha_{rt})g(n - 1) & \text{otherwise,} \end{cases}$$

$$\frac{\partial g(n)}{\partial g(n-1)} = \begin{cases} 1 - \alpha_{at} & \text{if } \hat{g}(n) < g(n - 1), \\ 1 - \alpha_{rt} & \text{otherwise,} \end{cases}$$



Yu, Chin-Yun, and György Fazekas. "Sound Matching an Analogue Levelling Amplifier Using the Newton-Raphson Method." AES AIMLA. 2025.

Chin-Yun Yu, Christopher Mitcheltree, Alistair Carson, Stefan Bilbao, Joshua D. Reiss, and György Fazekas, "Differentiable all-pole filters for time-varying audio systems," in DAFX, 2024, pp. 345–352.

Forward-mode AD

Newton's method needs second-order information

$$\theta \mapsto \theta - \left(\frac{\partial^2 \mathcal{L}}{\partial^2 \theta} \right)^{-1} \frac{\partial \mathcal{L}}{\partial \theta}$$

- The Hessian = the Jacobian of backprop = vector-to-vector function
- Jacobian-vector product (JVP) 

$$\text{JVP}_{\mathbf{J}} : \mathbf{v} \mapsto \mathbf{J}\mathbf{v}$$

reverse-mode AD (backprop)

$$\frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial b} \frac{\partial b}{\partial \theta} = \left(\text{VJP}_{\frac{\partial b}{\partial \theta}} \circ \text{VJP}_{\frac{\partial y}{\partial b}} \circ \text{VJP}_{\frac{\partial \mathcal{L}}{\partial y}} \right) (1)$$

$$= \left(\text{JVP}_{\frac{\partial \mathcal{L}}{\partial y}} \circ \text{JVP}_{\frac{\partial y}{\partial b}} \circ \text{JVP}_{\frac{\partial b}{\partial \theta}} \right) (1)$$

forward-mode AD

JVP of a scalar recursion

... is the same filter forward in time.

$$y(n) = x(n) + \alpha(n)y(n - 1)$$
$$\frac{\partial y(n)}{\partial \theta} = \frac{\partial x(n)}{\partial \theta} + \frac{\partial a(n)}{\partial \theta}y(n - 1) + \alpha(n)\frac{\partial y(n - 1)}{\partial \theta}$$

Yu, Chin-Yun, and György Fazekas. "Sound Matching an Analogue Levelling Amplifier Using the Newton-Raphson Method." AES AIMLA. 2025.

Different Hessian compute strategies

Table 1: The memory cost (in MB) and runtime (in milliseconds) of different hessian computation strategies in PyTorch.

	autograd rev-rev	func			
		rev-rev	fwd-rev	rev-fwd	fwd-fwd
Memory	1066	2358	3534	6306	5364
Time	26.5	28.1	28.1	64.7	38.9

Yu, Chin-Yun, and György Fazekas. "Sound Matching an Analogue Levelling Amplifier Using the Newton-Raphson Method." AES AIMLA. 2025.

Differentiable continuous delay length

Learning continuous delay time of a digital comb filter

$$y(n) = x(n) + \alpha y(n - M)$$

Frequency sampling: learning a sinusoidal

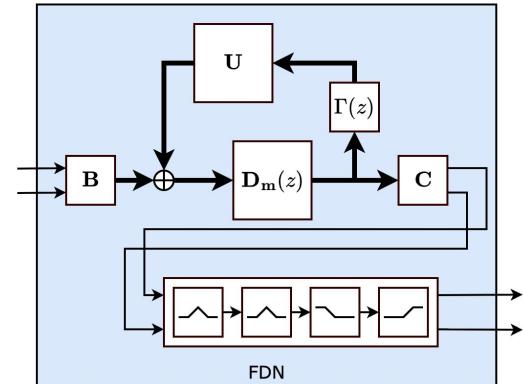
$$H(e^{j\omega}) = \frac{1}{1 - \alpha e^{-j\omega M}}$$

Time domain learning: non-trivial

- Ideal continuous time filter: **anti-causal, infinite order**
- Proposed: work on the poles directly

Fast AD of Feedback Delay Network

- Artificial reverberation
- Fast time domain training: possible
- Challenges
 - super high order (>10000) thus non-trivial to parallelise
 - differentiable delay lines



$$\begin{bmatrix} v_1(n + d_1) \\ v_2(n + d_2) \\ \vdots \\ v_M(n + d_M) \end{bmatrix} = \mathbf{A}\mathbf{v}(n) + \mathbf{B}x(n)$$

Parameter-varying stability

- LTI systems: coefficient parametrisation is enough
 - e.g., Biquad triangle
- LPV systems: LTI stable \neq LPV stable
 - Forms that guarantee stability
 - **Parallel filters**
 - Lattice filters

Gradients' numerical robustness

- TDF-II → more robust in forward pass
- DF-II → more robust in backpropagation
- Can we have the best of both worlds?
- Suggestion
 - TDF-II: forward-mode AD
 - DF-II reverse-mode AD
- Needs some benchmarks...

Special thanks

- Christopher Mitcheltree, Ben Hayes, and Rodrigo Diaz, for helpful feedback and comments from our conversations.
- Julius Smith for his valuable online materials about filters.

Resources

- 1) Yu, C.Y., Mitcheltree, C., Carson, A., Bilbao, S., Reiss, J.D. and Fazekas, G., 2024. Differentiable all-pole filters for time-varying audio systems. DAFx 2024
- 2) Yu, C.-Y., Fazekas, G., 2024. Differentiable Time-Varying Linear Prediction in the Context of End-to-End Analysis-by-Synthesis. Proc. Interspeech 2024, 1820-1824
- 3) Yu, C.-Y. and Fazekas, G., 2024 'GOLF: A Singing Voice Synthesiser with Glottal Flow Wavetables and LPC Filters', TISMIR, 7(1)
- 4) Yu, C.-Y. and Fazekas, G., 2025. "Sound Matching an Analogue Levelling Amplifier Using the Newton-Raphson Method." AES AIMLA 2025
- 5) Yu, C.-Y., . "DiffVox: A differentiable model for capturing and analysing vocal effects distributions." Proc. DAFx. 2025.
- 6) Yu, Chin-Yun, and György Fazekas. "Accelerating Automatic Differentiation of Direct Form Digital Filters." DlffSys Workshop @EurIPS 2025.
- 7) Yu, Chin-Yun. "PhilTorch: Accelerating Automatic Differentiation of Digital Filters In PyTorch". ADC 2025
- 8) Yu, C.-Y., 2025. Notes on differentiable TDF-II filter.
<https://iamycy.github.io/posts/2025/04/differentiable-tdf-ii/>
- 9) Smith, J.O. Introduction to Digital Filters with Audio Applications, <http://ccrma.stanford.edu/~jos/filters/>, online book, 2007 edition